



LINEAR CLASSIFIERS IN PYTHON

Welcome to the course!

Michael (Mike) Gelbart

Instructor

The University of British Columbia



Assumed knowledge

In this course we'll assume you have some prior exposure to:

- Python, at the level of *Intermediate Python for Data Science*
- scikit-learn, at the level of *Supervised Learning with scikit-learn*
- supervised learning, at the level of *Supervised Learning with scikit-learn*

Fitting and predicting

```
In [1]: import sklearn.datasets
```

```
In [2]: newsgroups = sklearn.datasets.fetch_20newsgroups_vectorized()
```

```
In [3]: X, y = newsgroups.data, newsgroups.target
```

```
In [4]: X.shape
```

```
Out[4]: (11314, 130107)
```

```
In [5]: y.shape
```

```
Out[5]: (11314,)
```

```
In [6]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [7]: knn = KNeighborsClassifier(n_neighbors=1)
```

```
In [8]: knn.fit(X,y)
```

```
In [9]: y_pred = knn.predict(X)
```



Model evaluation

```
In [10]: knn.score(X,y)
```

```
Out[10]: 0.99991
```

```
In [11]: from sklearn.model_selection import train_test_split
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
In [13]: knn.fit(X_train, y_train)
```

```
In [14]: knn.score(X_test, y_test)
```

```
Out[14]: 0.66242
```



LINEAR CLASSIFIERS IN PYTHON

Let's practice!



LINEAR CLASSIFIERS IN PYTHON

Applying logistic regression and SVM

Michael (Mike) Gelbart

Instructor

The University of British Columbia



Using LogisticRegression

```
In [1]: from sklearn.linear_model import LogisticRegression
```

```
In [2]: lr = LogisticRegression()
```

```
In [3]: lr.fit(X_train, y_train)
```

```
In [4]: lr.predict(X_test)
```

```
In [5]: lr.score(X_test, y_test)
```

LogisticRegression example

```
In [1]: import sklearn.datasets

In [2]: wine = sklearn.datasets.load_wine()

In [3]: from sklearn.linear_model import LogisticRegression

In [4]: lr = LogisticRegression()

In [5]: lr.fit(wine.data, wine.target)

In [6]: lr.score(wine.data, wine.target)
Out[6]: 0.972

In [7]: lr.predict_proba(wine.data[:1])
Out[7]: array([[ 9.951e-01,  4.357e-03,  5.339e-04]])
```


Using LinearSVC

LinearSVC works the same way:

```
In [1]: import sklearn.datasets  
  
In [2]: wine = sklearn.datasets.load_wine()  
  
In [3]: from sklearn.svm import LinearSVC  
  
In [4]: svm = LinearSVC()  
  
In [5]: svm.fit(wine.data, wine.target)  
  
In [6]: svm.score(wine.data, wine.target)  
Out[6]: 0.893
```

Using SVC

```
In [1]: import sklearn.datasets
In [2]: wine = sklearn.datasets.load_wine()
In [3]: from sklearn.svm import SVC
In [4]: svm = SVC() # default hyperparameters
In [5]: svm.fit(wine.data, wine.target);
In [6]: svm.score(wine.data, wine.target)
Out[6]: 1.
```

Model complexity review:

- **Underfitting:** model is too simple, low training accuracy
- **Overfitting:** model is too complex, low test accuracy



LINEAR CLASSIFIERS IN PYTHON

Let's practice!



LINEAR CLASSIFIERS IN PYTHON



Linear decision boundaries

Michael (Mike) Gelbart

Instructor

The University of British Columbia

Linear decision boundaries






Definitions

Vocabulary:


- **classification**: learning to predict categories
- **decision boundary**: the surface separating different predicted classes
- **linear classifier**: a classifier that learns linear decision boundaries
 - e.g., logistic regression, linear SVM
- **linearly separable**: a data set can be perfectly explained by a linear classifier

Linearly separable data

not linearly separable



linearly separable





LINEAR CLASSIFIERS IN PYTHON

Let's practice!